

**2 a. Which variable(s) control the number of spheres? **

Answer :SPHERES

2 b. Which variable(s) control the color(s) of the spheres?

Answer: r,g,b

2 c. Which variable(s) control the location of the spheres?

Answer: x,y,z

2 d. Which variable(s) control the size of each sphere?

Answer: radius

3) a)



```
#include <cuda_runtime_api.h>
#include "device_launch_parameters.h"
#include "J:\ami\ami\common\book.h"
#include <cuda_runtime_api.h>
#include "device_launch_parameters.h"
#include <stdio.h>
#include <cuda.h>
#include <stdlib.h>
#include <time.h>
#include "..\..\ami\ami\common\cpu_bitmap.h"
```

```

#include<math.h>

#define DIM 1024

#define rnd( x ) (x * rand() / RAND_MAX)
#define INF 2e10f

struct Sphere {
    float  r, b, g;
    float  radius;
    float  x, y, z;
    __device__ float hit(float ox, float oy, float *n) {
        float dx = ox - x;
        float dy = oy - y;
        if (dx*dx + dy*dy < radius*radius) {
            float dz = sqrtf(radius*radius - dx*dx - dy*dy);
            *n = dz / sqrtf(radius * radius);
            return dz + z;
        }
        return -INF;
    }
};

#define SPHERES 10

__constant__ Sphere s[SPHERES];

__global__ void kernel(unsigned char *ptr) {
    // map from threadIdx/BlockIdx to pixel position
    int x = threadIdx.x + blockIdx.x * blockDim.x;
    int y = threadIdx.y + blockIdx.y * blockDim.y;
    int offset = x + y * blockDim.x * gridDim.x;
    float  ox = (x - DIM / 2);
    float  oy = (y - DIM / 2);

    float  r = 0, g = 0, b = 0;
    float  maxz = -INF;
    for (int i = 0; i<SPHERES; i++) {
        float  n;
        float  t = s[i].hit(ox, oy, &n);
        if (t > maxz) {
            float fscale = n;
            r = s[i].r * fscale;
            g = s[i].g * fscale;

```

```

        b = s[i].b * fscale;
        maxz = t;
    }
}

ptr[offset * 4 + 0] = (int)(r * 255);
ptr[offset * 4 + 1] = (int)(g * 255);
ptr[offset * 4 + 2] = (int)(b * 255);
ptr[offset * 4 + 3] = 255;
}

// globals needed by the update routine
struct DataBlock {
    unsigned char *dev_bitmap;
};

int main(void) {
    DataBlock data;
    // seed the random function
    srand(time(NULL));
    // capture the start time
    cudaEvent_t start, stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);
    cudaEventRecord(start, 0);

    CPUBitmap bitmap(DIM, DIM, &data);
    unsigned char *dev_bitmap;

    // allocate memory on the GPU for the output bitmap
    cudaMalloc((void**)&dev_bitmap,
        bitmap.image_size());

    // allocate temp memory, initialize it, copy to constant
    // memory on the GPU, then free our temp memory
    Sphere *temp_s = (Sphere*)malloc(sizeof(Sphere)* SPHERES);
    for (int i = 0; i<SPHERES; i++) {
        temp_s[i].r = 0;
        temp_s[i].g = 0;
        temp_s[i].b = 1;
        temp_s[i].x = rnd(i * 100 - 500);
        temp_s[i].y = rnd(i * 100 - 500);
        temp_s[i].z = rnd(i * 100 - 500);
    }
}

```

```

        temp_s[i].radius = i * 10 + 20;
    }
    cudaMemcpyToSymbol(s, temp_s,
        sizeof(Sphere)* SPHERES);
    free(temp_s);

    // generate a bitmap from our sphere data
    dim3  grids(DIM / 16, DIM / 16);
    dim3  threads(16, 16);
    kernel << <grids, threads >> >(dev_bitmap);

    // copy our bitmap back from the GPU for display
    cudaMemcpy(bitmap.get_ptr(), dev_bitmap,
        bitmap.image_size(),
        cudaMemcpyDeviceToHost);

    // get stop time, and display the timing results
    cudaEventRecord(stop, 0);
    cudaEventSynchronize(stop);
    float  elapsedTime;
    cudaEventElapsedTime(&elapsedTime,
        start, stop);
    printf("Time to generate: %3.1f ms\n", elapsedTime);

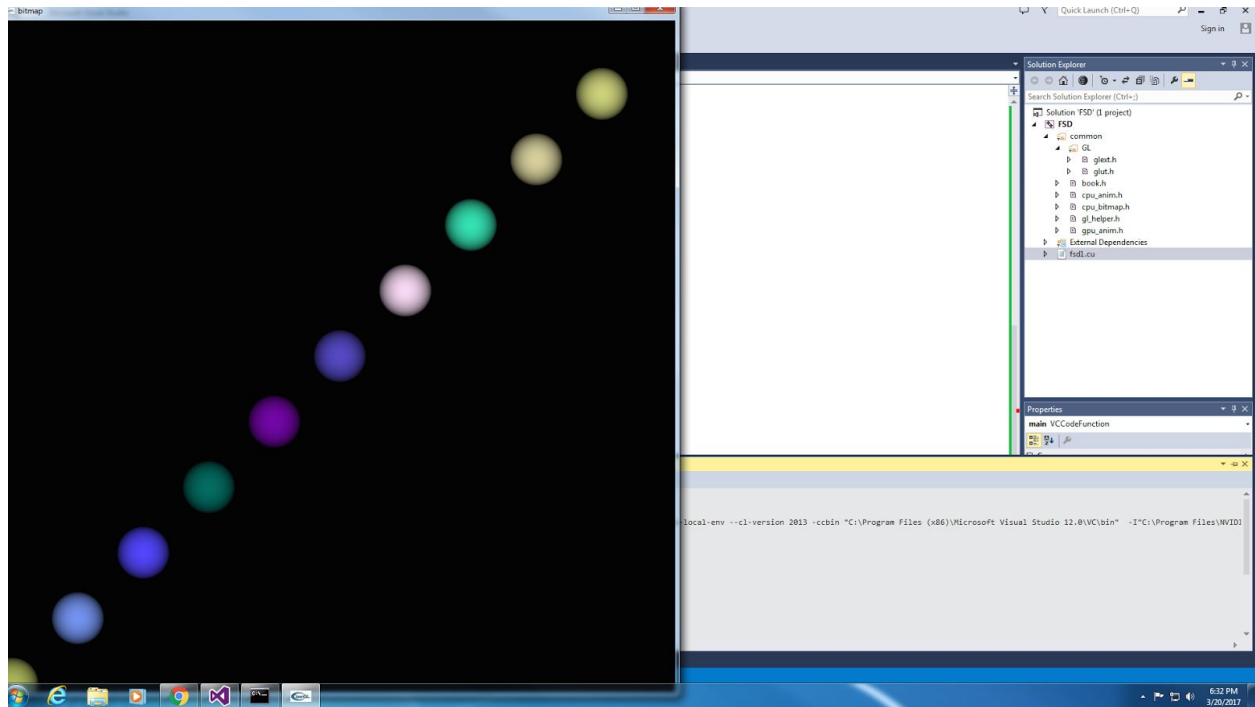
    cudaEventDestroy(start);
    cudaEventDestroy(stop);

    cudaFree(dev_bitmap);

    // display
    bitmap.display_and_exit();
}

```

3)b)



```
#include <cuda_runtime_api.h>
#include "device_launch_parameters.h"
#include "J:\\ami\\ami\\common\\book.h"
#include <cuda_runtime_api.h>
#include "device_launch_parameters.h"
#include <stdio.h>
#include <cuda.h>
#include <stdlib.h>
#include <time.h>
#include "..\\..\\ami\\ami\\common\\cpu_bitmap.h"
#include <math.h>

#define DIM 1024

#define rnd( x ) (x * rand() / RAND_MAX)
#define INF 2e10f

struct Sphere {
    float r, b, g;
    float radius;
    float x, y, z;
    __device__ float hit(float ox, float oy, float *n) {
```

```

        float dx = ox - x;
        float dy = oy - y;
        if (dx*dx + dy*dy < radius*radius) {
            float dz = sqrtf(radius*radius - dx*dx - dy*dy);
            *n = dz / sqrtf(radius * radius);
            return dz + z;
        }
        return -INF;
    }
};

#define SPHERES 10

__constant__ Sphere s[SPHERES];

__global__ void kernel(unsigned char *ptr) {
    // map from threadIdx/BlockIdx to pixel position
    int x = threadIdx.x + blockIdx.x * blockDim.x;
    int y = threadIdx.y + blockIdx.y * blockDim.y;
    int offset = x + y * blockDim.x * gridDim.x;
    float  ox = (x - DIM / 2);
    float  oy = (y - DIM / 2);

    float  r = 0, g = 0, b = 0;
    float  maxz = -INF;
    for (int i = 0; i < SPHERES; i++) {
        float  n;
        float  t = s[i].hit(ox, oy, &n);
        if (t > maxz) {
            float fscale = n;
            r = s[i].r * fscale;
            g = s[i].g * fscale;
            b = s[i].b * fscale;
            maxz = t;
        }
    }

    ptr[offset * 4 + 0] = (int)(r * 255);
    ptr[offset * 4 + 1] = (int)(g * 255);
    ptr[offset * 4 + 2] = (int)(b * 255);
    ptr[offset * 4 + 3] = 255;
}

```

// globals needed by the update routine

```

struct DataBlock {
    unsigned char *dev_bitmap;
};

int main(void) {
    DataBlock data;
    // seed the random function
    srand(time(NULL));
    // capture the start time
    cudaEvent_t start, stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);
    cudaEventRecord(start, 0);

    CPUBitmap bitmap(DIM, DIM, &data);
    unsigned char *dev_bitmap;

    // allocate memory on the GPU for the output bitmap
    cudaMalloc((void**)&dev_bitmap,
               bitmap.image_size());

    // allocate temp memory, initialize it, copy to constant
    // memory on the GPU, then free our temp memory
    Sphere *temp_s = (Sphere*)malloc(sizeof(Sphere)* SPHERES);
    for (int i = 0; i<SPHERES; i++) {
        temp_s[i].r = rnd(1.0f);
        temp_s[i].g = rnd(1.0f);
        temp_s[i].b = rnd(1.0f);
        temp_s[i].x = i * 100 - 500;
        temp_s[i].y = i * 100 - 500;
        temp_s[i].z = i * 100 - 500;
        temp_s[i].radius = 40;
    }
    cudaMemcpyToSymbol(s, temp_s,
                       sizeof(Sphere)* SPHERES);
    free(temp_s);

    // generate a bitmap from our sphere data
    dim3 grids(DIM / 16, DIM / 16);
    dim3 threads(16, 16);
    kernel << <grids, threads >> >(dev_bitmap);

    // copy our bitmap back from the GPU for display

```

```

        cudaMemcpy(bitmap.get_ptr(), dev_bitmap,
                    bitmap.image_size(),
                    cudaMemcpyDeviceToHost);

// get stop time, and display the timing results
cudaEventRecord(stop, 0);
cudaEventSynchronize(stop);
float elapsedTime;
cudaEventElapsedTime(&elapsedTime,
                    start, stop);
printf("Time to generate: %.3f ms\n", elapsedTime);


cudaEventDestroy(start);
cudaEventDestroy(stop);

cudaFree(dev_bitmap);

// display
bitmap.display_and_exit();
}

```

3) c)



```
#include <cuda_runtime_api.h>
```



```

#include "device_launch_parameters.h"
#include "J:\\ami\\ami\\common\\book.h"
#include <cuda_runtime_api.h>
#include "device_launch_parameters.h"
#include <stdio.h>
#include <cuda.h>
#include <stdlib.h>
#include <time.h>
#include "..\\ami\\ami\\common\\cpu_bitmap.h"
#include <math.h>

#define DIM 1024

#define rnd( x ) (x * rand() / RAND_MAX)
#define INF 2e10f

struct Sphere {
    float  r, b, g;
    float  radius;
    float  x, y, z;
    __device__ float hit(float ox, float oy, float *n) {
        float dx = ox - x;
        float dy = oy - y;
        if (dx*dx + dy*dy < radius*radius) {
            float dz = sqrtf(radius*radius - dx*dx - dy*dy);
            *n = dz / sqrtf(radius * radius);
            return dz + z;
        }
        return -INF;
    }
};

#define SPHERES 10

__constant__ Sphere s[SPHERES];

__global__ void kernel(unsigned char *ptr) {
    // map from threadIdx/BlockIdx to pixel position
    int x = threadIdx.x + blockIdx.x * blockDim.x;
    int y = threadIdx.y + blockIdx.y * blockDim.y;
    int offset = x + y * blockDim.x * gridDim.x;
    float  ox = (x - DIM / 2);
    float  oy = (y - DIM / 2);

```

```

float  r = 0, g = 0, b = 0;
float  maxz = -INF;
for (int i = 0; i<SPHERES; i++) {
    float  n;
    float  t = s[i].hit(ox, oy, &n);
    if (t > maxz) {
        float fscale = n;
        r = s[i].r * fscale;
        g = s[i].g * fscale;
        b = s[i].b * fscale;
        maxz = t;
    }
}

ptr[offset * 4 + 0] = (int)(r * 255);
ptr[offset * 4 + 1] = (int)(g * 255);
ptr[offset * 4 + 2] = (int)(b * 255);
ptr[offset * 4 + 3] = 255;
}

// globals needed by the update routine
struct DataBlock {
    unsigned char  *dev_bitmap;
};

int main(void) {
    DataBlock  data;
    // seed the random function
    srand(time(NULL));
    // capture the start time
    cudaEvent_t  start, stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);
    cudaEventRecord(start, 0);

    CPUBitmap bitmap(DIM, DIM, &data);
    unsigned char  *dev_bitmap;

    // allocate memory on the GPU for the output bitmap
    cudaMalloc((void**)&dev_bitmap,
        bitmap.image_size());

    // allocate temp memory, initialize it, copy to constant

```

```

// memory on the GPU, then free our temp memory
Sphere *temp_s = (Sphere*)malloc(sizeof(Sphere)* SPHERES);
for (int i = 0; i<SPHERES; i++) {
    temp_s[i].r = rnd(1.0f);
    temp_s[i].g = rnd(1.0f);
    temp_s[i].b = rnd(1.0f);
    temp_s[i].x = i * 100 - 500;
    temp_s[i].y = i * 100 - 500;
    temp_s[i].z = i * 100 - 500;
    temp_s[i].radius = rnd(80);
}
cudaMemcpyToSymbol(s, temp_s,
    sizeof(Sphere)* SPHERES);
free(temp_s);

// generate a bitmap from our sphere data
dim3  grids(DIM / 16, DIM / 16);
dim3  threads(16, 16);
kernel << <grids, threads >> >(dev_bitmap);

// copy our bitmap back from the GPU for display
cudaMemcpy(bitmap.get_ptr(), dev_bitmap,
    bitmap.image_size(),
    cudaMemcpyDeviceToHost);

// get stop time, and display the timing results
cudaEventRecord(stop, 0);
cudaEventSynchronize(stop);
float  elapsedTime;
cudaEventElapsedTime(&elapsedTime,
    start, stop);
printf("Time to generate:  %3.1f ms\n", elapsedTime);

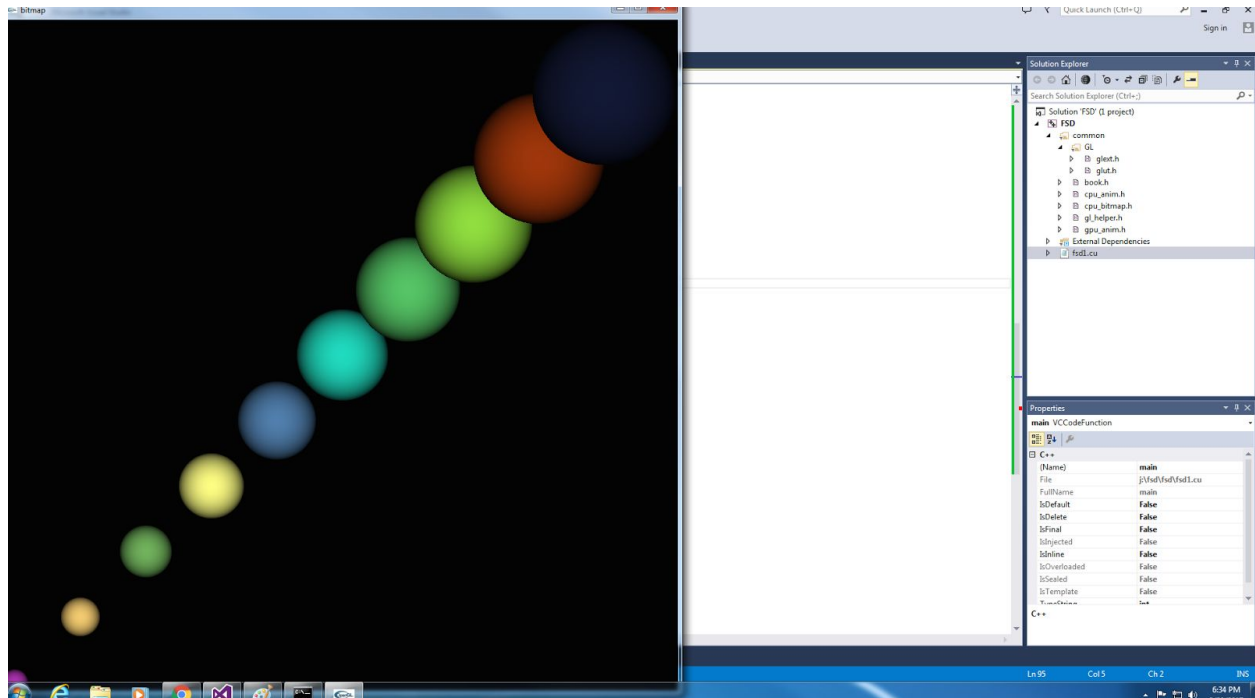
cudaEventDestroy(start);
cudaEventDestroy(stop);

cudaFree(dev_bitmap);

// display
bitmap.display_and_exit();
}

```

3)d)



```
#include <cuda_runtime_api.h>
#include "device_launch_parameters.h"
#include "J:\\ami\\ami\\common\\book.h"
#include <cuda_runtime_api.h>
#include "device_launch_parameters.h"
#include <stdio.h>
#include <cuda.h>
#include <stdlib.h>
#include <time.h>
#include "..\\..\\ami\\ami\\common\\cpu_bitmap.h"
#include <math.h>
```

```
#define DIM 1024
```

```
#define rnd( x ) (x * rand() / RAND_MAX)
```

```
#define INF 2e10f
```

```
struct Sphere {
    float  r, b, g;
    float  radius;
    float  x, y, z;
    __device__ float hit(float ox, float oy, float *n) {
```

```

        float dx = ox - x;
        float dy = oy - y;
        if (dx*dx + dy*dy < radius*radius) {
            float dz = sqrtf(radius*radius - dx*dx - dy*dy);
            *n = dz / sqrtf(radius * radius);
            return dz + z;
        }
        return -INF;
    }
};

#define SPHERES 10

__constant__ Sphere s[SPHERES];

__global__ void kernel(unsigned char *ptr) {
    // map from threadIdx/BlockIdx to pixel position
    int x = threadIdx.x + blockIdx.x * blockDim.x;
    int y = threadIdx.y + blockIdx.y * blockDim.y;
    int offset = x + y * blockDim.x * gridDim.x;
    float  ox = (x - DIM / 2);
    float  oy = (y - DIM / 2);

    float  r = 0, g = 0, b = 0;
    float  maxz = -INF;
    for (int i = 0; i < SPHERES; i++) {
        float  n;
        float  t = s[i].hit(ox, oy, &n);
        if (t > maxz) {
            float fscale = n;
            r = s[i].r * fscale;
            g = s[i].g * fscale;
            b = s[i].b * fscale;
            maxz = t;
        }
    }

    ptr[offset * 4 + 0] = (int)(r * 255);
    ptr[offset * 4 + 1] = (int)(g * 255);
    ptr[offset * 4 + 2] = (int)(b * 255);
    ptr[offset * 4 + 3] = 255;
}

```

// globals needed by the update routine

```

struct DataBlock {
    unsigned char *dev_bitmap;
};

int main(void) {
    DataBlock data;
    // seed the random function
    srand(time(NULL));
    // capture the start time
    cudaEvent_t start, stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);
    cudaEventRecord(start, 0);

    CPUBitmap bitmap(DIM, DIM, &data);
    unsigned char *dev_bitmap;

    // allocate memory on the GPU for the output bitmap
    cudaMalloc((void**)&dev_bitmap,
        bitmap.image_size());

    // allocate temp memory, initialize it, copy to constant
    // memory on the GPU, then free our temp memory
    Sphere *temp_s = (Sphere*)malloc(sizeof(Sphere)* SPHERES);
    for (int i = 0; i<SPHERES; i++) {
        temp_s[i].r = rnd(1.0f);
        temp_s[i].g = rnd(1.0f);
        temp_s[i].b = rnd(1.0f);
        temp_s[i].x = i * 100 - 500;
        temp_s[i].y = i * 100 - 500;
        temp_s[i].z = i * 100 - 500;
        temp_s[i].radius = i * 10 + 20;
    }
    cudaMemcpyToSymbol(s, temp_s,
        sizeof(Sphere)* SPHERES);
    free(temp_s);

    // generate a bitmap from our sphere data
    dim3 grids(DIM / 16, DIM / 16);
    dim3 threads(16, 16);
    kernel << <grids, threads >> >(dev_bitmap);

    // copy our bitmap back from the GPU for display

```

```
    cudaMemcpy(bitmap.get_ptr(), dev_bitmap,
               bitmap.image_size(),
               cudaMemcpyDeviceToHost);

    // get stop time, and display the timing results
    cudaEventRecord(stop, 0);
    cudaEventSynchronize(stop);
    float elapsedTime;
    cudaEventElapsedTime(&elapsedTime,
                       start, stop);
    printf("Time to generate: %3.1f ms\n", elapsedTime);

    cudaEventDestroy(start);
    cudaEventDestroy(stop);

    cudaFree(dev_bitmap);

    // display
    bitmap.display_and_exit();
}
```